

Training WG Updates for Kubeflow v1.5

Yuan Tang @terrytangyuan on behalf of Kubeflow Training WG



Major Changes

- Support Elastic Torch Training. #1453
- Migrated v1 MPI operator to the unified operator. #1457
- Python SDK is now on PyPI. https://pypi.org/project/kubeflow-training/
- Standalone XGBoost/PyTorch/MXNet operator repos have been archived. <u>kubeflow/community#479</u>



Bug Fixes and Improvements

- Added init container config to avoid DNS lookup failure for PyTorchJob. #1493
- Added E2E tests debugging guide. #1469
- Added missing replica type and index labels in TFJob. #1474
- Fixed issue when setting scheduler. #1448
- Fixed TFJob status when enableDynamicWorker is set to true. #1455



Roadmap

- Release training-operator v1.4 to be included in Kubeflow v1.5 release.
- Migrate v2 MPI operator to unified operator.
- Migrate PaddlePaddle operator to unified operator.
- Support elastic training for additional frameworks besides PyTorch.
- Support different gang scheduling definitions.
- Improve test coverage.

The most up-to-date roadmap can be found in this document.



Training WG Community

- Become a User Group Contributor for Training WG.
- Check the <u>Meeting Notes</u> to attend WG meetings and get the latest updates.
- Join our Slack channel: #wg-training.
- If you are using Training Operator, please update the <u>Adopters</u> list.
- How to contribute?
 - Check the <u>Developer Guide</u>.
 - Issues with the "Help Wanted" label.
 - Submit a new <u>Issue</u> with the new feature.





Thank you

https://github.com/kubeflow/training-operator

Yuan Tang @terrytangyuan